



Write a small paragraph about what you observed experimentally and what was your analytical/theoretical analysis of time complexity.

I observed that the first test case (array size of 5) always had a significantly higher runtime than other cases with a similar number of elements (array size of 10). This is odd, but may be due to the nature of the input file: many of the trials in this case have B arrays of size 1, which triggers my program to run a different, possibly slower, version of its iteration than other cases would. I also observed that the linear trend upwards became more stable at high array size; that is, there was a more predictable curve (indicated in orange on the plot) when the array size was greater. This is consistent with the fact that since these cases had a greater number of elements, they represent the general trend, rather than a local trend. The time complexity plot from experimental values appears to be somewhat linear, looking more similar to an $O(n)$ time complexity than the $O(n^2)$ time complexity I predicted. This makes sense because my $O(n^2)$ prediction was a worst case runtime analysis, but because of how my method was written the worst case runtime would very rarely occur: there is a nested for-loop triggered by an if-else statement that would always be in a worst case scenario, but the if-else is so specific that it usually wouldn't occur. Thus, only one for-loop is looped through, giving an $O(n)$ runtime in most cases, as I observed in the experimental plot.